

HealthAI: Intelligent Healthcare Assistant

"HealthAI: Intelligent Healthcare Assistant Using IBM Granite"

By:

Team ID :

LTVIP2025TMID31708

Team Size : 1

Team Leader : Guttula
Niranjan

Date: [26 06 2025]

TABLE OF CONTENTS

This section will be automatically generated. It will provide a comprehensive overview, listing all major headings and subheadings within this HealthAI project documentation.

1. INTRODUCTION

1.1 Project Overview

"HealthAI: Intelligent Healthcare Assistant Using IBM Granite" represents a pioneering endeavor aimed at revolutionizing personal healthcare management through cutting-edge artificial intelligence. This innovative project harnesses the formidable power of IBM Watson Machine Learning and advanced Generative AI capabilities to deliver truly intelligent and personalized healthcare assistance. The platform's strategic design incorporates several integral features: a 'Patient Chat' module for intuitive, real-time conversational support for health-related inquiries; 'Disease Prediction' employing sophisticated machine learning models to identify potential health conditions based on user-provided symptoms and data; 'Treatment Plans' that generate tailored, evidence-based recommendations; and 'Health Analytics' providing users with actionable insights derived from their health data. Currently, HealthAI is in its critical ideation, design, and planning phases, establishing a robust framework for its future development and deployment.

1.2 Purpose

The core objective of this comprehensive documentation is to furnish a detailed and structured blueprint for the HealthAI project. It meticulously chronicles every aspect from the initial conceptualization and problem identification through the intricate design, meticulous planning, and foundational testing strategies. This document serves as an indispensable, centralized repository of information, designed to ensure clarity and alignment among all involved stakeholders, including project team members, technical architects, quality assurance

professionals, business strategists, and potential investors. Furthermore, it precisely outlines the critical healthcare challenges HealthAI seeks to address, elucidates the innovative AI-driven solutions, details the proposed scalable and secure architectural design, enumerates all functional and non-functional requirements, and presents a pragmatic, phased project plan. Its existence is crucial for guiding development, facilitating informed decision-making, and ensuring the successful realization of the HealthAI vision.

2. IDEATION PHASE

The Ideation Phase is crucial for defining the core problems the HealthAI platform aims to solve, understanding the needs and perspectives of the target users, and generating a broad range of potential solutions before refining and prioritizing the most promising ideas. This phase leverages collaborative brainstorming and user empathy techniques to ensure the resulting solution is both impactful and aligned with real-world needs.

2.1 Problem Statements

Identifying clear, user-centric problem statements is the foundation of the ideation process. These statements capture the challenges faced by potential users in their own context and guide the development of relevant solutions. Below are 2-3 key problem statements derived from initial research and discussions, structured to highlight the user, their goal, the obstacle, its root cause, and the emotional impact.

Problem Statement 1: Navigating Health Information Overload

- **I am (Customer):** A busy professional managing multiple responsibilities.
- **I'm trying to:** Quickly find trustworthy and relevant health information when I experience a symptom or have a health question without spending hours searching online or waiting for a doctor's appointment.
- **But:** The sheer volume of conflicting and often unreliable health information available online makes it difficult to discern credible sources and understand what is applicable to my specific situation.
- **Because:** There is no single, reliable, easily accessible platform that provides personalized, evidence-based health information and guidance tailored to my symptoms and health profile.
- **Which makes me feel:** Anxious, confused, and uncertain about how to address my health concerns effectively and efficiently.

Problem Statement 2: Managing Chronic Conditions Effectively

- **I am (Customer):** A senior citizen living with multiple chronic conditions (e.g., diabetes, hypertension).
- **I'm trying to:** Consistently track my health metrics (blood sugar, blood pressure), manage my medication schedule, adhere to treatment plans, and communicate effectively with my healthcare providers.
- **But:** It is challenging to manually keep track of various readings, remember multiple medication timings, understand complex treatment instructions, and consolidate all this information for doctor visits.
- **Because:** I lack an integrated, user-friendly tool that simplifies the management and monitoring of my multiple health needs and facilitates easy sharing of information with my care team.
- **Which makes me feel:** Overwhelmed, isolated, and worried about my ability to manage my health independently and prevent complications.

Problem Statement 3: Accessing Timely and Personalized Health Guidance in Underserved Areas

- **I am (Customer):** A person living in a rural or medically underserved area.
- **I'm trying to:** Access timely, personalized, and potentially life-saving health information and initial guidance based on my symptoms without the need for immediate physical travel to a distant healthcare facility.
- **But:** Geographic distance, limited access to local healthcare professionals, and long waiting times for appointments make it difficult to get quick assessments or advice for non-emergency health issues.
- **Because:** There are insufficient accessible, AI-powered virtual health tools that can provide reliable symptom analysis and preliminary guidance, leveraging relevant local health context where possible.
- **Which makes me feel:** Vulnerable, neglected, and apprehensive about potential health issues going unaddressed.

2.2 Empathy Map Canvas

To deeply understand the target users of HealthAI, an Empathy Map Canvas was developed. This tool helps visualize the user's perspective by exploring what they Say, Think, Do, and Feel in relation to managing their health and interacting with healthcare systems. This provides valuable insights into their pain points, needs, and motivations, guiding the design of a user-centric solution. The following captures the perspective of a typical user seeking quick, reliable health information.

Says:

- "I wish there was an easy way to figure out what these symptoms mean without booking a doctor's appointment."
- "Online health forums are confusing; everyone says something different."
- "My doctor is great, but getting an appointment takes forever for a simple question."
- "Is this a normal ache, or something I should worry about?"
- "I need to remember to take my pills, but I keep forgetting."
- "Finding reliable diet advice is hard."

Thinks:

- 'Could this symptom be serious?'
- 'Am I overreacting, or should I get this checked?'
- 'I hope this isn't going to cost a lot of money.'
- 'What if the information I find online is wrong?'
- 'It would be so much easier if I had all my health info in one place.'
- 'How can I explain all my symptoms clearly?'

Does:

- Searches symptoms on Google or WebMD.
- Asks friends or family for medical advice.
- Delays seeking professional medical help for minor issues.
- Tries over-the-counter remedies or home cures.
- Manually tracks symptoms or vital signs using notebooks or separate apps.
- Reads health articles and blogs.

- Forgets to take medications occasionally.

Feels:

- Anxious when uncertain about symptoms.
- Frustrated by difficulty finding reliable information.
- Overwhelmed by health management tasks.
- Relieved when receiving clear, actionable advice.
- Worried about potential misdiagnosis or overlooking a serious condition.
- Hopeful when a potential solution seems promising.
- Isolated when dealing with health issues alone.

2.3 Brainstorming & Idea Prioritization

Following the definition of problem statements and gaining deeper user empathy, the team engaged in a brainstorming session to generate potential features and functionalities for HealthAI. The goal was to explore a wide array of possibilities before evaluating their potential impact and feasibility.

Step 1: Team Collaboration & Problem Selection

The core project team gathered for a dedicated session. Utilizing the previously defined problem statements and the insights from the Empathy Map, the team collectively reviewed and affirmed the primary challenges HealthAI would address: the need for accessible, reliable health information; the difficulties in managing chronic conditions; and the gap in timely guidance, particularly in underserved areas. These problems served as the focal points for the subsequent idea generation.

Step 2: Brainstorming, Idea Listing, and Grouping

A free-form brainstorming session encouraged the team to propose any feature or concept that could potentially alleviate the identified user pain points. Ideas were captured without immediate judgment and subsequently grouped into logical categories to create structure and identify related functionalities.

- **Core AI Functions**
 - AI Symptom Checker (using NLP on user input)
 - Disease Prediction based on symptoms and user profile
 - Personalized Treatment Plan Generator
 - AI-Powered Patient Chatbot (conversational health assistant)
- **User Management & Engagement**
 - User Profile & Health Data Management (secure storage)
 - Health Analytics Dashboard (visualizing trends)
 - Medication Reminder System
 - Appointment & Follow-up Tracker
 - Personalized Health Goal Setting & Tracking
- **Integration & Utilities**
 - Integration with Wearable Devices (for data sync)
 - Integration with Pharmacy Services (prescription refills)
 - Telehealth Integration (linking to virtual doctor visits)

- Health Content Library (curated articles, videos)

Step 3: Idea Prioritization

The brainstormed features were then evaluated based on two key criteria: their potential **Impact** on solving the identified user problems and achieving project goals, and their technical **Feasibility** within the project's scope, resources, and timeline, especially considering the use of IBM Granite and Watson services. A simple prioritization matrix (High, Medium, Low) was used, with High Impact/High Feasibility features receiving the highest priority.

The priority scoring logic was applied as follows:

- High Impact & High Feasibility = **Priority 1 (High)**
- High Impact & Medium Feasibility = **Priority 2 (Medium-High)**
- Medium Impact & High Feasibility = **Priority 2 (Medium-High)**
- High Impact & Low Feasibility = **Priority 3 (Medium-Low)**
- Medium Impact & Medium Feasibility = **Priority 3 (Medium-Low)**
- Low Impact (regardless of Feasibility) = **Priority 4 (Low)**
- Medium Impact & Low Feasibility = **Priority 4 (Low)**
- Low Impact & Any Feasibility = **Priority 4 (Low)**

The prioritization results are presented in the table below:

Feature	Category	Impact	Feasibility	Priority Score/Level
AI Symptom Checker	Core AI Functions	High	High	1 (High)
AI-Powered Patient Chatbot	Core AI Functions	High	High	1 (High)
User Profile & Health Data Management	User Management & Engagement	High	High	1 (High)
Personalized Treatment Plan Generator	Core AI Functions	High	Medium	2 (Medium-High)
Disease Prediction	Core AI Functions	High	Medium	2 (Medium-High)
Health Analytics Dashboard	User Management & Engagement	Medium	High	2 (Medium-High)
Medication Reminder System	User Management	Medium	High	2 (Medium-High)

Feature	Category	Impact	Feasibility	Priority Score/Level
	& Engagement			
Personalized Health Goal Setting & Tracking	User Management & Engagement	Medium	Medium	3 (Medium-Low)
Health Content Library	Integration & Utilities	Medium	Medium	3 (Medium-Low)
Appointment & Follow-up Tracker	User Management & Engagement	Medium	Medium	3 (Medium-Low)
Integration with Wearable Devices	Integration & Utilities	Medium	Low	4 (Low)
Integration with Pharmacy Services	Integration & Utilities	Medium	Low	4 (Low)
Telehealth Integration	Integration & Utilities	Low	Low	4 (Low)

This prioritization helped the team focus on the most impactful and feasible features for the initial phases of development, forming the basis for the subsequent requirements definition and planning.

3. REQUIREMENT ANALYSIS

The Requirement Analysis phase is pivotal in translating the project's vision and user needs into a clear, actionable set of specifications. This section delves into understanding the user's journey, defining the system's capabilities and constraints, mapping data flows, capturing user perspectives through stories, and outlining the foundational technology stack that will power HealthAI. By thoroughly analyzing these aspects, the project team establishes a solid blueprint for design and development, ensuring the final product effectively addresses user pain points and delivers on its promise of intelligent healthcare assistance.

3.1 Customer Journey Map

Understanding the user's experience from initial awareness of a health concern through interaction with HealthAI and subsequent follow-up is crucial for designing an intuitive and effective platform. The Customer Journey Map below illustrates the typical path a user might take, highlighting their actions, thoughts, emotions, pain points in the traditional process, and the opportunities for HealthAI to provide a superior experience.

Stage 1: Symptom Recognition / Initial Concern

- **User Actions:** Experiences a new or worsening symptom (e.g., headache, fatigue, unusual pain). Feels generally unwell or notices a specific change in their health.
- **User Thoughts:** 'What is this symptom?', 'Is this serious?', 'Should I be worried?', 'Is this something I can treat myself or do I need a doctor?'.
- **User Emotions:** Concern, slight anxiety, curiosity, confusion.
- **Pain Points (Traditional):** Uncertainty about the symptom's cause or severity. Lack of immediate access to medical opinion.
- **Opportunities for HealthAI:** Provide an initial, low-barrier point of access for symptom assessment and information seeking. Offer reassurance and context.

Stage 2: Information Seeking (Pre-HealthAI)

- **User Actions:** Searches online (Google, WebMD, health forums) for symptom information. Reads articles, forums, blogs. Asks friends or family for advice.
- **User Thoughts:** 'There's so much information, where do I start?', 'Is this information reliable?', 'This contradicts what I read elsewhere.', 'This seems too serious/too mild.', 'What worked for my friend might not work for me.'
- **User Emotions:** Frustration, confusion, increased anxiety due to conflicting or alarming information, skepticism.
- **Pain Points (Traditional):** Information overload, unreliability of sources, difficulty personalizing general information, lack of structured guidance.
- **Opportunities for HealthAI:** Act as a single, trusted source for information. Filter and personalize information based on user input. Provide structured, evidence-based guidance.

Stage 3: Considering Professional Help

- **User Actions:** Considers calling a doctor's office. Checks insurance coverage. Looks up local clinics. Decides whether the symptom warrants an appointment.
- **User Thoughts:** 'Can I get an appointment soon?', 'Will my insurance cover this?', 'Is it worth taking time off work/spending money for this?', 'Maybe it will just go away.', 'Should I go to urgent care?'
- **User Emotions:** Hesitation, financial concern, inconvenience, apprehension about potential diagnosis, hope for resolution.
- **Pain Points (Traditional):** Difficulty getting timely appointments, cost concerns, uncertainty about the necessity of a visit, logistical hurdles.
- **Opportunities for HealthAI:** Help users assess the urgency of their situation. Offer preliminary insights that inform the decision to seek professional help. Suggest when professional consultation is highly recommended. Integrate with telehealth or appointment booking systems (future).

Stage 4: HealthAI Interaction (Initiation)

- **User Actions:** Discovers HealthAI (e.g., through search, recommendation, ad). Downloads the app or visits the website. Registers or logs in. Initiates interaction (e.g., clicks "Check Symptoms" or opens the chat).
- **User Thoughts:** 'Let's see if this can help.', 'Is this easy to use?', 'Will it understand me?', 'Is my data safe?', 'How accurate is this AI?'
- **User Emotions:** Hope, curiosity, slight skepticism, caution regarding data privacy.

- **Pain Points (Traditional):** (This is the point HealthAI starts to alleviate traditional pain points). The pain point here is the initial barrier of trusting and learning a new tool.
- **Opportunities for HealthAI:** Provide a smooth onboarding experience. Clearly communicate value proposition and data security measures. Ensure intuitive interface design.

Stage 5: HealthAI Core Interaction (Symptom Analysis & Chat)

- **User Actions:** Enters symptoms via chat or structured form. Answers clarifying questions from HealthAI. Provides additional relevant health information (age, gender, existing conditions, etc., based on profile or input). Uses the AI chatbot for conversational queries.
- **User Thoughts:** 'This is much faster than searching.', 'The questions are helpful.', 'It seems to understand what I'm saying.', 'Is this analysis reliable?', 'How is it using my health profile?', 'Can I ask follow-up questions?'
- **User Emotions:** Engagement, relief as information is structured, growing trust (if results are relevant), frustration (if system misunderstands), comfort (with conversational chat).
- **Pain Points (Pre-HealthAI Interaction):** Difficulty structuring symptom information, inability to get personalized follow-up questions, receiving generic, non-interactive information.
- **Opportunities for HealthAI:** Leverage IBM Granite for natural, empathetic chat. Use IBM Watson Machine Learning for accurate symptom analysis and context understanding. Provide clear, easy-to-understand responses. Continuously learn and improve interactions.

Stage 6: Receiving Insights & Recommendations

- **User Actions:** Views HealthAI's output: potential conditions, contributing factors, personalized information, suggested next steps (e.g., monitor, self-care, see a doctor soon, seek urgent care), generated treatment plan recommendations (if applicable), health analytics insights.
- **User Thoughts:** 'Okay, this makes sense.', 'It's considering my [condition/age].', 'This recommendation is clear.', 'It's suggesting I see a doctor, that's helpful.', 'These treatment steps seem manageable.', 'My health trend is visible here.'
- **User Emotions:** Clarity, relief, validation, empowerment (with actionable steps), potential concern (if serious conditions are flagged), understanding.
- **Pain Points (Pre-HealthAI Interaction):** Receiving overwhelming or non-actionable information, lack of personalized context, no clear next steps, inability to see personal health trends easily.
- **Opportunities for HealthAI:** Present complex information simply. Clearly differentiate between AI insights and medical advice. Provide evidence-based recommendations tailored to the user profile and input. Integrate different insights (symptoms, history, analytics) into a cohesive view.

Stage 7: Taking Action / Follow-up

- **User Actions:** Follows HealthAI's recommendations (e.g., applies self-care steps, monitors symptoms, schedules doctor's appointment, adheres to treatment plan, uses medication reminders, tracks health metrics). Updates their health profile or logs new data in HealthAI. Uses HealthAI for ongoing management.

- **User Thoughts:** 'I'll try these steps.', 'I need to remember to schedule that appointment.', 'HealthAI can remind me about my medication.', 'Logging this data helps me track progress.', 'How is my health changing over time?'
- **User Emotions:** Proactiveness, discipline, hope for improvement, sense of control, satisfaction (if actions lead to positive outcomes).
- **Pain Points (Pre-HealthAI Interaction):** Difficulty remembering medication schedules, inconsistent tracking of health data, feeling alone in managing health, lack of tools to support adherence to plans.
- **Opportunities for HealthAI:** Provide proactive reminders and notifications. Facilitate easy data logging and tracking. Offer tools for goal setting and progress monitoring. Support ongoing engagement and health management.

3.2 Solution Requirements

Defining the requirements is fundamental to building a system that meets user needs and achieves project goals. This section details the functional capabilities and non-functional constraints that govern the design and implementation of the HealthAI platform.

Functional Requirements

Functional requirements specify what the system must do. They are grouped into epics, representing major areas of functionality, with supporting user stories or sub-tasks.

Epic: User Management

- User Registration: Allow new users to create an account using email/password or third-party authentication.
 - Sub-task: Email verification upon registration.
 - Sub-task: Secure password handling (hashing).
- User Login: Enable registered users to securely log in to their account.
- Profile Management: Allow users to view and update their personal information (age, gender, location, pre-existing conditions, allergies, medications).
 - Sub-task: Data validation for profile fields.
 - Sub-task: Option to add/remove health conditions, allergies, medications.
- Password Reset/Recovery: Provide a mechanism for users to reset their password if forgotten.
- Account Deactivation/Deletion: Allow users to deactivate or permanently delete their account and associated data.
- Role Management (Admin): Enable administrators to manage user accounts (e.g., view, suspend, delete - future scope).

Epic: Symptom Analysis & Disease Prediction

- Symptom Input: Allow users to input symptoms via natural language text or a structured selection interface.
 - Sub-task: Process natural language symptom descriptions using NLP.
 - Sub-task: Provide suggested symptoms based on user input.
- Clarifying Questions: The system must ask relevant follow-up questions based on initial symptom input to narrow down potential conditions.
 - Sub-task: Adapt questions based on user's profile data (age, gender, history).
 - Sub-task: Handle ambiguous or incomplete user responses.

- Potential Condition Identification: Analyze symptoms and user profile data to identify potential underlying health conditions.
 - Sub-task: Rank potential conditions by likelihood.
 - Sub-task: Display supporting symptoms or factors for each potential condition.
- Risk Assessment: Provide a general risk assessment (e.g., low, medium, high) based on the analysis.
- Recommendation for Next Steps: Based on the analysis, recommend appropriate next steps (e.g., self-care, monitor symptoms, consult doctor, seek urgent care, call emergency services).
 - Sub-task: Clearly state that recommendations are not a substitute for professional medical advice.

Epic: Chatbot Interaction (Patient Chat)

- Conversational Interface: Provide an intuitive chat interface for users to ask health-related questions in natural language.
 - Sub-task: Understand a wide range of health queries.
 - Sub-task: Maintain context throughout a conversation.
- Information Provision: Provide clear, concise, and evidence-based answers to user questions, leveraging the HealthAI knowledge base and AI capabilities.
 - Sub-task: Cite sources or explain the basis for the information where appropriate.
 - Sub-task: Handle out-of-scope or inappropriate queries gracefully.
- Guided Conversations: Guide users through symptom analysis or information gathering processes via chat.
- Escalation/Redirection: Identify when a query requires professional medical attention and recommend consulting a doctor.

Epic: Treatment Plans & Recommendations

- Generate Treatment Recommendations: Based on identified conditions or user queries, generate personalized recommendations for self-care, lifestyle changes, or potential treatment approaches.
 - Sub-task: Recommendations should align with established medical guidelines where possible.
 - Sub-task: Consider user's profile (allergies, existing medications) when generating recommendations.
- Present Treatment Steps: Break down complex recommendations into actionable steps.
- Medication Tracking & Reminders: Allow users to add their medications and schedule reminders.
 - Sub-task: Configure reminder frequency and time.
 - Sub-task: Mark doses as taken.
- Treatment Adherence Tracking: Allow users to track their adherence to recommended steps or medication schedules.

Epic: Health Data Analytics & Tracking

- Data Input: Allow users to log various health data points (e.g., blood pressure, blood sugar, weight, mood, activity levels).
 - Sub-task: Support manual data entry.

- Sub-task: (Future) Integrate with wearable devices or other health apps for automated data syncing.
- Data Visualization: Present logged health data in clear, interactive charts and graphs.
 - Sub-task: Show trends over time (daily, weekly, monthly).
 - Sub-task: Compare current data points against historical averages or target ranges.
- Personalized Insights: Provide insights based on the user's tracked data, identifying patterns or potential areas of concern.
 - Sub-task: Alert users to significant deviations or trends based on their profile or set goals.
- Report Generation: (Future) Generate shareable reports of health data for users to share with their healthcare providers.

Epic: Knowledge Base & Content Management

- Curated Health Library: Provide access to a library of reliable health articles, FAQs, and information on common conditions.
 - Sub-task: Content should be easily searchable and browsable.
 - Sub-task: Content should be reviewed for accuracy and clarity.
- Content Management System (Admin): (Future) Provide an interface for administrators to add, edit, and publish health content.

Non-functional Requirements

Non-functional requirements specify criteria that can be used to judge the operation of a system, rather than specific behaviors. They define the system's quality attributes.

- **Usability:**

The system must be easy to learn and use for individuals of varying technical proficiency, including older adults. The interface should be intuitive, with clear navigation and minimal steps required to access core functionalities like symptom checking or chatbot interaction. Error messages should be helpful and guiding. Onboarding should be simple and quick.

- **Security:**

Protecting sensitive user health data is paramount.

- All data transmitted between the user interface and the backend must be encrypted (e.g., using TLS/SSL).
- Sensitive data stored in the database (e.g., health records) must be encrypted at rest.
- User authentication and authorization must be strictly enforced (e.g., using industry-standard protocols and IBM Cloud IAM).
- API keys and credentials for accessing external services (IBM Granite, Watson) must be managed securely.
- The system should be protected against common web vulnerabilities (e.g., XSS, SQL Injection).
- Regular security audits and penetration testing should be planned.

- Adherence to relevant health data privacy regulations (e.g., HIPAA, GDPR) is mandatory.
- **Reliability:**

The system should perform its intended functions correctly and consistently.

 - Minimize downtime and system errors.
 - Implement robust error handling and logging mechanisms.
 - Ensure data consistency across the application.
 - Provide regular data backups and a disaster recovery plan.
- **Performance:**

The system should respond quickly to user requests.

 - Symptom analysis and chatbot responses should be returned within a few seconds (e.g., < 5 seconds for core interactions).
 - Page load times should be optimized (e.g., < 3 seconds).
 - The system should be able to handle a significant number of concurrent users without degradation in performance.
 - Database queries should be optimized for speed.
- **Availability:**

The system should be accessible to users whenever they need it.

 - Target system uptime of 99.5% or higher.
 - Deploy applications across multiple availability zones to protect against single points of failure.
 - Implement load balancing to distribute traffic.
 - Have strategies for handling planned and unplanned outages.
- **Scalability:**

The system must be able to handle a growing number of users and increasing data volume.

 - The architecture should support horizontal scaling of application servers and database read replicas.
 - Utilize cloud services that offer elastic scaling capabilities.
 - Design databases and data models with future growth in mind.
 - Ensure ML model inference can scale with request volume.

3.3 Data Flow Diagrams & User Stories

Visualizing how data moves through the system (DFD) and defining the system's features from the user's perspective (User Stories) are essential steps in requirement analysis.

Data Flow Diagrams

A Level 0 Data Flow Diagram (Context Diagram) provides a high-level view of the system, illustrating HealthAI as a single process interacting with external entities and accessing primary data stores. This diagram shows the boundary of the system and the major data flows in and out.

The HealthAI system (the central process) interacts with primary external entities: Users (Patients) and potentially External Healthcare Systems (future integration for data exchange or referrals). Users provide input (symptoms, health data, queries) and receive output (analysis results, treatment plans, information, dashboard views). The system accesses and updates internal data stores: the User Profiles Database (storing personal and health history), the Health Records Database (storing logged health data and interaction history), and the Knowledge Base (containing curated health information and model data). The system also interacts with External APIs (like IBM Granite and Watson services) to process requests and generate AI-driven insights, representing another external interaction point though often conceptualized as a service rather than a primary external entity on a Level 0 diagram focused on users and data stores.

[PLACEHOLDER FOR LEVEL 0 DATA FLOW DIAGRAM - Illustrating data flow between external entities, main processes, and data stores]

User Stories

User stories capture the desired functionality from the perspective of different users, following the "As a [User Type], I want to [Action], So that [Benefit]" format. They include acceptance criteria, priority, and release information to guide development.

User Story	Acceptance Criteria	Priority	Release
As a Mobile User , I want to input my symptoms via text or voice, so that I can get a preliminary analysis without typing extensively.	<ul style="list-style-type: none"> Given I am on the symptom check screen, when I tap the input field, I can choose text or voice input. When I provide symptoms via text or voice, the system processes the input and displays potential conditions and next steps. The system 	High	Sprint 1

User Story	Acceptance Criteria	Priority	Release
	<ul style="list-style-type: none"> asks clarifying questions if initial input is insufficient. 		
As a Web User , I want to view a summary of my health data trends on a dashboard after logging in, so that I can quickly see my progress and identify patterns.	<ul style="list-style-type: none"> Given I am a registered user, when I successfully log in, I am redirected to the user dashboard. The dashboard displays key health metrics (e.g., BP, blood sugar) with graphical trends for a default period (e.g., last 7 days). I can select different time periods (e.g., 30 days, 90 days) to view trends. 	High	Sprint 2
As a Mobile User , I want to receive push	<ul style="list-style-type: none"> Given I have added a 	High	Sprint 2

User Story	Acceptance Criteria	Priority	Release
notifications for my scheduled medication reminders, so that I don't forget to take my medicine.	<p>medication with a scheduled time, when the scheduled time arrives and I have not marked it as taken, I receive a push notification on my device.</p> <ul style="list-style-type: none"> • The notification includes the medication name and dosage. • I can dismiss the notification or tap it to open the app to the medication tracking screen. 		
As an Admin User , I want to add a new health article to the curated content library, so that users have access to up-to-date health information.	<ul style="list-style-type: none"> • Given I am logged in as an administrator, when I navigate to the content management section, I can 	Medium	Future Release

User Story	Acceptance Criteria	Priority	Release
	<p>initiate adding a new article.</p> <ul style="list-style-type: none"> I can input the article title, body text, relevant tags, and source information. Upon saving, the article is added to the content library and is accessible to users. 		
As a Web User , I want to be able to securely update my pre-existing conditions and allergies in my health profile, so that HealthAI can provide more personalized recommendations .	<ul style="list-style-type: none"> Given I am logged in, when I navigate to my profile settings, I can view my current health profile details. I can add new conditions or allergies from a predefined list or by searching. I can remove existing 	High	Sprint 1

User Story	Acceptance Criteria	Priority	Release
	conditions or allergies. <ul style="list-style-type: none"> • Changes are saved and reflected in my profile upon confirmation. 		

3.4 Technology Stack

The technology stack forms the foundation upon which the HealthAI platform is built. Careful consideration of each component ensures the system is robust, scalable, secure, and leverages the power of IBM's AI and cloud services effectively.

Technical Architecture

The HealthAI architecture follows a layered approach, typically including a presentation layer (User Interface), an application layer (Backend Logic), a data layer (Database), and external service integrations (External APIs, ML Models). This structure promotes separation of concerns, modularity, and scalability.

[PLACEHOLDER FOR TECHNOLOGY STACK ARCHITECTURE DIAGRAM - Illustrating the interplay of UI, backend, database, APIs, and ML components]

- **User Interface:**

The user interface will be developed using modern web and mobile frameworks to ensure a responsive and intuitive user experience.

- **Web:** React.js - A popular JavaScript library for building dynamic and performant single-page applications. Provides component-based architecture for reusability and maintainability.
- **Mobile:** React Native - Allows building native mobile applications for iOS and Android from a single codebase, accelerating development and ensuring consistency across platforms.
- **Styling/Markup:** HTML5, CSS3, possibly a CSS framework like Bootstrap or Material UI for rapid styling and responsiveness.

- **Application Logic (Backend):**

The backend will handle business logic, data processing, user authentication, and interaction with AI/ML services and the database.

- **Framework:** Python with Flask or Django - Python is well-suited for AI/ML integrations. Flask is lightweight for API development, while Django offers a more

comprehensive framework including an ORM and admin panel (useful for future admin features). Choice will depend on complexity needs.

- **Runtime Environment:** Python 3.x.
- **Hosting:** Deployed on IBM Cloud Code Engine, Red Hat OpenShift on IBM Cloud, or IBM Cloud Kubernetes Service for containerized deployment and scalability. Serverless functions (IBM Cloud Functions) could be used for specific event-driven tasks (e.g., processing data updates).

- **Database:**

A robust and scalable database is needed to store user profiles, health data, interaction history, and the system's knowledge base.

- **Choice:** IBM Cloud Databases for PostgreSQL or IBM Cloud Databases for MongoDB. PostgreSQL provides strong relational integrity for structured user profiles and historical data, while MongoDB (NoSQL) offers flexibility for potentially unstructured or semi-structured health log data. A combination might be considered, or one chosen based on the predominant data structure. These are managed cloud services offering built-in scalability, backups, and high availability.

- **External APIs:**

Integration with external services is critical for AI capabilities and potentially other functionalities.

- **Generative AI:** IBM Granite API (via IBM watsonx.ai) - Core API for leveraging large language models for conversational chat, generating treatment plan outlines, and potentially structuring symptom input.
- **Machine Learning:** IBM Watson Machine Learning API (via IBM watsonx.ai) - Used for deploying and running custom or pre-trained models for symptom analysis, disease prediction, and health data insights.
- **Authentication:** IBM App ID - Provides secure user authentication and identity management as a cloud service, supporting various identity providers.
- **Other Potential APIs:** Integration with third-party health data APIs (wearables - future), pharmacy APIs (future), or telehealth platforms (future).

- **Machine Learning Models:**

Specific ML models are central to HealthAI's intelligent functions.

- **Symptom Checker/Disease Prediction:** Classification models trained on medical datasets, potentially using frameworks like TensorFlow or PyTorch. These models would be deployed and managed via IBM Watson Machine Learning.
- **Natural Language Processing (NLP):** Leveraging IBM Watson Natural Language Understanding (NLU) services or custom models for processing user symptom descriptions and chat inputs, extracting medical entities and understanding intent.
- **Recommendation Engine:** Models for generating personalized treatment plan recommendations based on user profile, condition, and preferences.
- **Analytics Models:** Models for identifying patterns or anomalies in user health data.

Application Characteristics

The architectural choices and technology selections contribute significantly to the application's quality attributes.

- **Open-Source Frameworks:**

The project leverages widely adopted open-source technologies to benefit from community support, continuous improvement, and cost-effectiveness. Key frameworks include:

- **React.js/React Native:** For front-end development, enabling rapid development and a large developer community.
- **Python (Flask/Django):** A powerful and versatile language with extensive libraries for backend development and data science.
- **TensorFlow/PyTorch/Scikit-learn:** Standard libraries for building and training machine learning models.
- **Docker/Kubernetes:** Essential for containerization and orchestration, facilitating consistent deployment environments and scalability.

- **Security Implementations:**

Security is integrated at multiple levels:

- **Data Encryption:** Utilizing built-in encryption features of IBM Cloud Databases for data at rest. Implementing TLS/SSL for data in transit.
- **Identity and Access Management (IAM):** Using IBM Cloud IAM roles and policies to control access to cloud resources and services (Databases, Watson APIs, etc.). IBM App ID for user authentication and managing user identities securely.
- **API Security:** Using API keys or token-based authentication (e.g., OAuth 2.0 via App ID) to secure API endpoints.
- **Network Security:** Configuring firewalls and security groups to restrict network access to necessary components.
- **Code Security:** Following secure coding practices and utilizing tools for vulnerability scanning in the CI/CD pipeline.

- **Scalable Architecture:**

The design facilitates scaling to accommodate a growing user base and data volume:

- **Microservices/Modular Design:** Breaking down functionalities (e.g., Symptom Analysis Service, Chatbot Service, User Service) allows scaling individual components based on demand.
- **Containerization (Docker):** Packaging applications in containers ensures they run consistently across different environments and are easily deployable on orchestration platforms.
- **Orchestration (Kubernetes/Code Engine):** IBM Cloud Kubernetes Service or Code Engine automatically manages the deployment, scaling, and operation of application containers, allowing horizontal scaling based on load.
- **Cloud Databases:** Managed database services like IBM Cloud Databases offer built-in scaling options and read replicas to handle increasing database load.

- **Managed AI Services:** IBM Watson and Granite APIs are managed services designed to scale automatically based on API request volume.
- **Availability:**

High availability is ensured through redundancy and cloud infrastructure features:

- **Multi-Zone Deployment:** Deploying application containers and database instances across multiple availability zones within IBM Cloud regions protects against data center failures.
- **Load Balancing:** Using load balancers to distribute incoming traffic across healthy instances of application components, ensuring requests are served even if one instance fails.
- **Managed Database Redundancy:** IBM Cloud Databases offer built-in high availability configurations with automatic failover to a replica in case of primary instance failure.
- **Backup and Recovery:** Implementing automated backup strategies for data and application configurations to enable quick recovery in case of data loss or corruption.
- **Performance:**

Strategies are in place to optimize response times and throughput:

- **Optimized Database Access:** Using efficient queries, proper indexing, and connection pooling.
- **Caching:** Implementing caching mechanisms (e.g., Redis on IBM Cloud) for frequently accessed data or AI model outputs to reduce database/API calls and latency.
- **Asynchronous Processing:** Handling non-blocking operations (e.g., sending email verification) asynchronously to free up request threads using message queues (e.g., Kafka on IBM Cloud).
- **Content Delivery Network (CDN):** Utilizing a CDN for serving static assets (images, CSS, JS) to improve load times for users globally.
- **Efficient AI Model Inference:** Optimizing ML models for faster prediction times and leveraging hardware acceleration provided by IBM Cloud (e.g., GPUs if needed).

4. PROJECT DESIGN

The Project Design phase translates the validated problems, user needs, and prioritized ideas from the Ideation and Requirement Analysis phases into a concrete plan for the HealthAI platform. This involves defining how the proposed solution fits the identified problems, detailing the solution itself, and outlining the technical architecture that will support it. This phase ensures that the technical implementation aligns with the strategic goals and user requirements.

4.1 Problem-Solution Fit

Ensuring Problem-Solution Fit is critical to developing a product that truly addresses customer needs and is adopted successfully. This analysis maps the identified customer pain points and behaviors to how the HealthAI solution directly resolves them and strategies for engaging users.

Problem

HealthAI is designed to address several key pain points experienced by individuals managing their health:

- **Information Overload and Unreliability:** Users struggle to find accurate, trustworthy, and personalized health information online amidst a vast sea of conflicting and often misleading content. This leads to confusion and anxiety.
- **Difficulty Managing Chronic Conditions:** Individuals with chronic illnesses find it challenging to consistently track multiple health metrics, adhere to complex medication schedules and treatment plans, and effectively communicate their status to healthcare providers.
- **Lack of Timely Access to Guidance:** Accessing quick, preliminary health guidance for non-emergency symptoms is difficult due to appointment waiting times, geographic distance (especially in underserved areas), or cost barriers, leaving users uncertain and delaying potential intervention.
- **Fragmented Health Data:** Personal health information (symptoms, lab results, medication history, tracked metrics) is often scattered across different apps, paper records, or recalled from memory, making holistic health management and communication with doctors inefficient.
- **Anxiety and Uncertainty:** Unaddressed health questions or symptoms, coupled with difficulties in navigating the healthcare system, contribute significantly to user stress and apprehension about their health.

Behavioral Patterns

Customers exhibit specific behaviors influenced by these problems:

- **Frequent Online Self-Diagnosis:** Relying heavily on search engines (like Google) to understand symptoms, often encountering alarming or irrelevant results.
- **Seeking Advice from Non-Professionals:** Consulting friends, family, or online forums for health advice when professional help is inconvenient or perceived as unnecessary for minor issues.
- **Delayed Consultation:** Postponing doctor visits for symptoms perceived as minor or due to logistical/cost barriers, potentially missing early signs of serious conditions.
- **Inconsistent Tracking:** Attempting to track health metrics or medication manually but failing to maintain consistency due to lack of an integrated, easy-to-use system.
- **Passive Information Consumption:** Reading general health articles but struggling to apply the information to their specific personal context.
- **Reluctance to Seek Professional Help for 'Minor' Issues:** Avoiding overburdening healthcare providers with questions that seem insignificant, even if they cause personal anxiety.

Solution Adoption

HealthAI is designed to seamlessly integrate with and improve upon existing user behaviors or introduce beneficial new ones, thereby promoting adoption:

- **Enhancing Information Seeking:** Instead of chaotic online searches, users are guided to input symptoms into a structured AI system or ask questions via a natural language

chatbot. This leverages the user's existing inclination to seek information while providing a reliable, curated, and personalized source.

- **Simplifying Chronic Condition Management:** HealthAI replaces manual tracking or fragmented apps with an integrated platform for logging metrics, setting medication reminders, and visualizing data, directly supporting and improving the user's efforts to manage their conditions.
 - **Providing Immediate Preliminary Guidance:** HealthAI offers instant access to symptom analysis and initial recommendations, fitting into the user's need for timely information without the barriers of traditional healthcare access. This empowers users to take informed next steps, whether self-care or deciding to seek professional help.
 - **Consolidating Health Data:** By providing a secure central repository for user profiles, logged data, and interaction history, HealthAI encourages users to maintain a comprehensive digital health record, replacing fragmented systems.
 - **Reducing Anxiety:** By providing clear, personalized, and reliable information and guidance, HealthAI directly addresses the emotional impact of health uncertainty and the difficulty in accessing help. The conversational nature of the chatbot (using IBM Granite) can provide a more empathetic interaction than static symptom checkers.
- Touchpoints

Engaging with potential and existing users across various channels (touchpoints) is essential for awareness, acquisition, and retention:

- **Mobile App Stores (Apple App Store, Google Play Store):** Primary channel for discovery and download. Focus on compelling descriptions, screenshots, and user reviews.
- **Web Portal:** Providing a web-based version of the platform offers accessibility from desktops and an alternative access point. Useful for initial discovery via search engines.
- **Search Engine Marketing/Optimization (SEM/SEO):** Optimizing HealthAI's online presence for relevant health-related search terms (e.g., "symptom checker AI", "chronic condition management app").
- **Social Media:** Running targeted campaigns on platforms like Facebook, Instagram, and Twitter, sharing health tips, feature highlights, and user testimonials. Building a community around health and wellness.
- **Content Marketing:** Publishing valuable health-related content (blog posts, articles) on the HealthAI website, driving organic traffic and establishing authority.
- **Partnerships:**
 - **Healthcare Providers:** Collaborating with clinics, hospitals, or doctor networks to potentially recommend HealthAI as a supplementary tool for patient education and data tracking.
 - **Insurance Companies:** Exploring partnerships to integrate HealthAI services or offer them as a benefit to policyholders.
 - **Employer Wellness Programs:** Offering HealthAI as part of corporate wellness initiatives.
 - **Wearable Device Manufacturers:** Joint marketing or deeper technical integration efforts (future).
- **Health & Wellness Communities:** Engaging with online and offline groups focused on specific conditions or general wellness.

- **Public Relations:** Securing media coverage in health, technology, and business publications.
- **User Feedback Channels:** Providing clear in-app and website channels for users to provide feedback, ask questions, and get support, fostering loyalty and continuous improvement.

4.2 Proposed Solution

The HealthAI platform is envisioned as an intelligent, accessible, and comprehensive digital healthcare assistant leveraging the power of IBM's advanced AI capabilities. It aims to empower individuals to better understand and manage their health.

Problem Statement

The primary problem HealthAI addresses is the significant challenge individuals face in accessing timely, reliable, and personalized health information and guidance, compounded by difficulties in managing their health proactively and effectively in an increasingly complex healthcare landscape. This leads to anxiety, confusion, delayed care, and suboptimal health outcomes.

Idea / Solution Description

HealthAI is a cross-platform (web and mobile) application designed to serve as a user's first line of intelligent health assistance. At its core, it utilizes IBM Watson Machine Learning models for sophisticated health data analysis and IBM Granite, a state-of-the-art generative AI model, for natural, empathetic conversational interactions.

- **Patient Chat:** A central feature where users can interact with an AI chatbot using natural language. Powered by IBM Granite, this chatbot can answer health-related questions, provide information on symptoms and conditions, explain medical terms, and offer general wellness advice in a conversational, user-friendly manner. It is designed to be supportive and understanding.
- **Disease Prediction:** Leveraging IBM Watson Machine Learning models trained on vast medical datasets, this feature allows users to input symptoms through a guided interface or chat. The system analyzes the input, considers the user's health profile (age, gender, history), asks clarifying questions, and provides a list of potential conditions ranked by likelihood, along with explanations. It clearly states that this is not a diagnosis but an informational tool.
- **Treatment Plans:** Based on identified potential conditions, user goals, or specific health queries, HealthAI can generate personalized recommendations for self-care, lifestyle adjustments, or steps to discuss with a healthcare professional. IBM Granite can assist in structuring these plans conversationally, making them easy to understand and follow. This feature also includes tools for tracking adherence and setting reminders (e.g., medication reminders).
- **Health Analytics:** Users can log various health metrics (blood pressure, blood sugar, weight, activity, etc.) manually or, in future iterations, via connected devices. HealthAI uses IBM Watson ML to analyze this data over time, presenting trends visually through dashboards and providing personalized insights and alerts based on significant changes or patterns.
- **User Profile & Data Management:** A secure section where users store and manage their personal health information, which is used by the AI models to personalize

interactions and insights. Data privacy and security are paramount, adhering to relevant regulations.

- **Curated Health Library:** Access to a library of reliable, curated articles and information vetted for accuracy, providing foundational knowledge on common health topics accessible outside of chat interactions.

Novelty / Uniqueness

While digital health tools exist, HealthAI distinguishes itself through several key aspects:

- **Integrated AI Power:** Uniquely combines the analytical strength of IBM Watson Machine Learning for diagnosis/prediction and data analysis with the advanced conversational capabilities of IBM Granite for truly natural, empathetic patient interaction. Most existing symptom checkers or chatbots are based on simpler decision trees or less sophisticated NLP models.
- **Conversational Empathy (via Granite):** Leveraging the generative power of IBM Granite allows for more human-like, context-aware, and supportive dialogue in the Patient Chat, enhancing user comfort and engagement significantly beyond typical rule-based chatbots.
- **Holistic Integration:** Offers a single platform integrating symptom analysis, potential condition insights, personalized recommendations/plans, and health data tracking, providing a more cohesive health management experience than using separate apps for each function.
- **Focus on Guidance, Not Just Information:** Moves beyond simply providing data to offering actionable, personalized next steps and management plans.
- **Scalability and Security:** Built on the robust, secure, and scalable IBM Cloud infrastructure, designed from the ground up to handle sensitive health data and growing user loads effectively.

Social Impact / Customer Satisfaction

HealthAI is designed to deliver significant positive social impact and high customer satisfaction:

- **Improved Health Literacy and Empowerment:** Provides users with accessible, understandable, and reliable information, empowering them to make more informed decisions about their health.
- **Reduced Healthcare Anxiety:** Offers a readily available resource for initial symptom checks and health questions, alleviating stress associated with uncertainty and waiting for appointments.
- **Enhanced Proactive Health Management:** Tools for tracking data and receiving personalized insights encourage users to take a more active role in managing chronic conditions and maintaining wellness.
- **Potential for Early Detection:** By making preliminary symptom analysis widely accessible, HealthAI could potentially help users identify potential issues sooner, leading to earlier consultation with professionals when needed.
- **Increased Accessibility:** Provides valuable health information and guidance, particularly benefiting individuals in rural or underserved areas with limited access to immediate healthcare consultations.
- **Support for Healthcare Professionals:** Users coming to appointments with structured symptom descriptions or tracked data from HealthAI can potentially facilitate more efficient consultations.

- **Customer Satisfaction:** A user-friendly interface, personalized insights, reliable information, and responsive AI are expected to lead to high user satisfaction, fostering trust and continued engagement.

Business Model

A potential business model for HealthAI could involve a tiered approach:

- **Freemium Model:**
 - **Free Tier:** Basic access to Patient Chat (with limits on complexity/length), limited symptom checks per month, basic health metric tracking (manual input only), and access to the standard health library.
 - **Premium Subscription Tier:** Unlocks unlimited Patient Chat sessions, advanced symptom analysis with higher detail, personalized treatment plan generation, full Health Analytics dashboard with advanced insights and reporting, medication reminders, potential for wearable device integration (future), and priority support. This tier could be offered monthly or annually.
- **B2B Partnerships:**
 - **Healthcare Providers/Hospitals:** Offer a white-labeled or integrated version of HealthAI as a patient engagement tool, potentially with features for sharing user data (with explicit user consent) or facilitating virtual consultations. Licensing fees or per-patient fees.
 - **Insurance Companies:** Partner to offer the premium tier as a value-add benefit to policyholders, potentially reducing overall healthcare costs through improved preventive care and condition management.
 - **Employers:** Offer access to employees as part of corporate wellness programs.
- **API Access (Future):** Offer access to the core AI models (Symptom Analysis, Treatment Plan generation) as APIs for other health tech companies to integrate into their own platforms.
- **Data Monetization (Aggregated & Anonymized):** With strict user consent and complete anonymization/aggregation, insights from population-level health data trends could potentially be licensed to research institutions or public health bodies (adhering strictly to privacy regulations). This would NOT involve selling individual user data.

Scalability

The HealthAI solution is designed with scalability as a fundamental principle:

- **Cloud-Native Architecture:** Built and deployed on IBM Cloud, leveraging managed services that are inherently scalable (e.g., IBM Cloud Code Engine/Kubernetes, IBM Cloud Databases, IBM Watson/watsonx services).
- **Microservices Design:** Breaking the application logic into smaller, independent services (User Service, Chat Service, Analytics Service, Prediction Service) allows each component to be scaled independently based on its specific load requirements.
- **Elastic Computing:** Utilizing auto-scaling features of the cloud platform ensures that compute resources (like Code Engine instances or Kubernetes pods) automatically increase during peak demand and decrease during low periods, optimizing cost and performance.

- **Scalable Databases:** Using managed database services (IBM Cloud Databases for PostgreSQL/MongoDB) that offer vertical (scaling up) and horizontal (read replicas, sharding) scaling options as data volume and query load grow.
- **Managed AI Services:** IBM Watson and Granite services are designed to handle varying levels of API requests, scaling automatically within the limits of the service plan.
- **Load Balancing:** Distributing incoming user traffic across multiple instances of application services ensures no single instance becomes a bottleneck.
- **Stateless Services:** Designing backend services to be largely stateless simplifies horizontal scaling, as any incoming request can be handled by any available instance.

4.3 Solution Architecture

The HealthAI solution architecture is designed to be modular, scalable, secure, and highly available, leveraging the strengths of IBM Cloud and its AI/ML services. It follows a standard multi-tier pattern.

High-Level Architecture

The high-level architecture illustrates the primary components and their interactions. Users interact with the system via the User Interface. This interface communicates with the Application Logic/Backend Services, which handle business rules, user authentication, data processing, and orchestration. The Backend Services interact with the Database for data storage and retrieval and communicate with External APIs, including IBM Granite and IBM Watson Machine Learning services, to perform AI-driven tasks. Machine Learning Models represent the specific deployed models accessed via the External APIs.

The key components are:

- **User Interface (UI):** The user-facing application layer.
- **Application Logic/Backend Services:** The core processing layer.
- **Database:** The persistent data storage layer.
- **External APIs:** Interfaces to external services, notably IBM AI services.
- **Machine Learning Models:** The AI assets deployed and utilized.

[PLACEHOLDER FOR SOLUTION ARCHITECTURE DIAGRAM - Illustrating high-level components (UI, Application Logic, Database, External APIs, ML Models) and data flow]

Component Descriptions and Roles

Each component plays a specific role in the HealthAI ecosystem:

- **User Interface (UI):**
Role: Provides the graphical interface through which users interact with HealthAI. This includes web browsers and native mobile applications (iOS, Android). It handles user input (text, voice, selections), displays information (chat responses, analysis results, dashboards), and manages the user session presentation. It communicates with the Application Logic via APIs.
- **Application Logic/Backend Services:**
Role: Acts as the central nervous system of the application. It contains the core business logic, handles user authentication and authorization (potentially via a dedicated service

like IBM App ID), processes user requests, interacts with the Database to retrieve and store data, and makes calls to External APIs (IBM Watson/Granite) for AI processing. This layer orchestrates the flow of data and execution of tasks based on user actions.

- **Database:**

Role: Securely stores all persistent data required by the application. This includes user profiles (demographics, health history, conditions, allergies, medications), logged health data (metrics, symptom entries, chat history), and potentially the curated content library. It provides reliable data storage, retrieval, and management capabilities.

- **External APIs:**

Role: Serve as integration points to leverage specialized services outside the core backend. Key APIs include:

- **IBM Granite API:** Provides access to the generative AI model for powering natural language conversations in the Patient Chat and potentially generating structured text for recommendations.
- **IBM Watson Machine Learning API:** Used to deploy and invoke custom or pre-trained ML models for tasks such as symptom analysis, disease likelihood calculation, health data trend analysis, and pattern identification.
- **IBM App ID API:** Manages user identity, authentication, and authorization.
- Other potential third-party APIs (e.g., for sending notifications, integrating with wearables - future).

- **Machine Learning Models:**

Role: These are the trained AI models responsible for the "intelligent" aspects of HealthAI. They are hosted and managed within the IBM Watson Machine Learning environment. Their role is to take specific inputs (e.g., symptom descriptions, health metrics) and produce outputs such as potential condition probabilities, personalized insights, or structured recommendations based on their training data and algorithms.

Key Characteristics

The architecture incorporates key characteristics to ensure a robust and reliable system:

- **Security:**

Security is paramount for a healthcare application handling sensitive patient data. Measures include:

- **Data Encryption:** Encryption of data both in transit (using TLS/SSL for all communication between UI, Backend, Database, and APIs) and at rest (utilizing database encryption features).
- **Identity and Access Management (IAM):** Implementing strong authentication for users (via IBM App ID) and fine-grained authorization controls (role-based access control) to ensure users can only access their own data and administrators have appropriate permissions. IAM policies on IBM Cloud control access to infrastructure resources.

- **Secure API Access:** Protecting API endpoints using authentication tokens or API keys, ensuring only authorized backend services or users can interact with them.
- **Network Security:** Configuring virtual private networks (VPNs) or private endpoints for backend-database communication and using security groups/firewalls to restrict access to infrastructure.
- **Compliance:** Designing the architecture and data handling processes with relevant healthcare data privacy regulations (e.g., HIPAA, GDPR) in mind, although full certification would be a separate project phase.
- **Auditing and Monitoring:** Implementing logging and monitoring to detect and respond to suspicious activity.

- **Scalability:**

The architecture is built to scale horizontally to accommodate growth in user base and data:

- **Containerization and Orchestration:** Using Docker containers managed by IBM Cloud Code Engine or Kubernetes allows for easy deployment and scaling of backend services by adding more container instances.
- **Stateless Backend Services:** Designing the backend services to be stateless ensures that new instances can be added or removed seamlessly without impacting ongoing user sessions (session state is managed client-side or in an external cache).
- **Managed Scalable Database:** Utilizing IBM Cloud Databases which offer automatic scaling or easy configuration of replicas and sharding to handle increased database load and data volume.
- **Managed AI Services:** IBM Watson/Granite scale automatically based on API request volume.

- **Availability:**

Ensuring the system is accessible when users need it is critical for a health application:

- **Multi-Zone Deployment:** Deploying application and database instances across multiple availability zones within an IBM Cloud region protects against outages affecting a single physical location.
- **Load Balancing:** Distributing incoming user requests across multiple healthy instances of backend services.
- **Database High Availability:** Using managed database services with built-in replication and automatic failover mechanisms.
- **Resilient Design:** Implementing retry mechanisms and circuit breakers in backend services to handle temporary failures in dependent services (like APIs or databases).
- **Automated Backups and Disaster Recovery:** Regular automated backups of database and application configurations, along with a defined disaster recovery plan.

- **Performance:**

Optimizing response times and throughput is essential for a good user experience, especially with AI interactions:

- **Optimized Database Queries:** Designing efficient database schemas and queries, utilizing indexing, and connection pooling.
- **Caching:** Implementing caching layers (e.g., using IBM Cloud Databases for Redis) for frequently accessed data (like parts of the health library or recent AI responses) to reduce the load on the database and AI services.
- **Efficient AI Model Inference:** Optimizing the deployed ML models for performance and leveraging the underlying infrastructure provided by IBM Cloud (potentially including GPU acceleration if needed for complex models).
- **Asynchronous Processing:** Using message queues (e.g., IBM MQ or Kafka on IBM Cloud) for tasks that don't require an immediate response (like sending notifications or processing offline data) to keep the main request threads free.
- **Content Delivery Network (CDN):** Utilizing a CDN for serving static assets of the UI to minimize latency for users globally.

5. PROJECT PLANNING & SCHEDULING

The project planning and scheduling phase translates the defined requirements and architectural design into a actionable roadmap for development. This section outlines the HealthAI project's agile plan, detailing the product backlog, sprint schedule, team velocity, and the visualization of progress through a burndown chart. This structured approach ensures transparency, adaptability, and efficient resource allocation throughout the development lifecycle.

5.1 Project Planning Template

Product Backlog

The Product Backlog for HealthAI is a prioritized, dynamic list of all known features, functionalities, enhancements, and bug fixes required for the platform. Each item is described from the user's perspective (user story), assigned a 'Story Points' value indicating its estimated complexity and effort, and a 'Priority' to guide development order. The initial backlog focuses on the high-impact, high-feasibility features identified during the Ideation Phase, ensuring early delivery of core value to the user.

User Story/Feature	Description	Story Points	Priority
User Profile & Data Management	As a user, I want to securely create, update, and manage my personal and health profile (e.g., age, gender, pre-existing conditions, allergies, medications), so that HealthAI can provide accurate and personalized	5	High

User Story/Feature	Description	Story Points	Priority
	recommendations .		
AI Symptom Checker	As a user, I want to input my symptoms via text or voice, and receive a preliminary analysis including potential conditions and recommended next steps, so that I can quickly understand my health situation.	8	High
AI-Powered Patient Chatbot	As a user, I want to converse naturally with an AI assistant for health-related questions, receiving clear and empathetic answers, so that I can get immediate, reliable information and guidance.	13	High
Health Analytics Dashboard	As a user, I want to view a graphical summary of my tracked health data trends (e.g., blood pressure, blood sugar) over time, so that I can monitor my progress and identify patterns in my health.	8	Medium
Medication Reminder System	As a user, I want to set up and receive push notifications for	3	Medium

User Story/Feature	Description	Story Points	Priority
	my scheduled medication reminders, so that I consistently adhere to my treatment plan and don't miss doses.		
Personalized Treatment Plan Generation	As a user, I want to receive tailored self-care and lifestyle recommendations , and suggested treatment approaches based on my identified conditions or health goals, so that I have actionable steps to improve my well-being.	13	High
Admin Content Management (Future)	As an admin, I want to manage and publish curated health articles and FAQs to the HealthAI knowledge base, so that users have access to up-to-date and reliable health information.	5	Low

Sprint Schedule

The development of HealthAI will follow an Agile Scrum methodology, broken down into fixed-length sprints. Each sprint will have a clear objective and a set of user stories committed for completion. The initial plan outlines two hypothetical 2-week sprints, with remaining items allocated to future releases or considered for the next phases of development. This iterative approach allows for flexibility, continuous feedback, and regular delivery of working software.

Sprint Number	Duration	Objective	Planned User Stories/Features
Sprint 1	2 Weeks	Establish foundational user management and initial symptom inquiry capabilities, enabling secure profile creation and basic health assessment.	<ul style="list-style-type: none"> User Profile & Data Management (5 Story Points) AI Symptom Checker (8 Story Points)
Sprint 2	2 Weeks	Implement core conversational AI for general health queries and basic health data visualization, enhancing user engagement and self-monitoring.	<ul style="list-style-type: none"> AI-Powered Patient Chatbot (13 Story Points) Health Analytics Dashboard (8 Story Points) Medication Reminder System (3 Story Points)
Future Sprints / Releases	Ongoing	Develop and refine advanced personalization, integration features, and administrative functionalities.	<ul style="list-style-type: none"> Personalized Treatment Plan Generation (13 Story Points) Admin Content Management (5 Story Points) <i>Other prioritized backlog items...</i>

Velocity

In Agile project management, velocity measures the amount of work a development team can complete in a single sprint, typically expressed in 'Story Points'. It is a key metric for forecasting and planning future sprints. For the HealthAI project, based on the completed work in Sprint 1 (13 story points) and planned work in Sprint 2 (24 story points), we assume an average team velocity of approximately **20 story points per sprint**. This average helps in predicting how much work can be realistically committed to and delivered in subsequent sprints, aiding in consistent progress and realistic expectation setting. This metric will be continuously refined and adjusted as the team progresses and gains a clearer understanding of its actual capacity and efficiency.

Burndown Chart

A burndown chart is a graphical representation of work left to do versus time. It visually tracks the progress of a sprint or an entire project, showing the remaining work (typically in story points) against the planned timeline. For HealthAI, a burndown chart would illustrate how the total accumulated story points are being 'burned down' (completed) over the duration of the development sprints. It helps the team and stakeholders quickly assess progress, identify potential bottlenecks, and understand if the project is on track to meet its completion goals. A typical burndown chart would show an ideal progress line (straight diagonal from total work to zero) and an actual progress line, allowing for easy comparison and immediate insight into project health.

[PLACEHOLDER FOR PROJECT BURNDOWN CHART - Showing remaining work vs. time over sprints]

6. FUNCTIONAL AND PERFORMANCE TESTING

The Functional and Performance Testing phase is critical to ensure that the HealthAI platform not only meets its specified requirements but also operates efficiently and reliably under various conditions. This section details a series of test scenarios designed to validate core functionalities and assess the system's performance characteristics, providing insights into its readiness for deployment. Each test case includes a unique identifier, a description of the scenario, the steps to execute, the expected outcome, a simulated actual result, and a pass/fail status.

Functional Test Cases

These test cases verify that specific features and functionalities of HealthAI operate as intended, meeting the defined functional requirements.

Test Case ID	Scenario	Test Steps	Expected Result	Actual Result	Pass/Fail
TC_F_001	User Registration - Invalid Email Format	1. Navigate to the user registration page.	System displays an error message indicating invalid email format and	System displayed "Please enter a valid email address." error message,	Pass

Test Case ID	Scenario	Test Steps	Expected Result	Actual Result	Pass/Fail
		<ol style="list-style-type: none"> Enter a username and a password. Input an invalid email address (e.g., "user@domain"). Attempt to complete registration. 	prevents registration.	preventing registration.	
TC_F_002	Patient Chat - Basic Conversational Query	<ol style="list-style-type: none"> Log in as a registered user. Access the Patient Chat interface. Type "I have a sore throat and cough. What should I do?" Observe the chatbot's response. 	Chatbot provides initial self-care advice for common cold/flu symptoms and asks for clarifying details (e.g., fever presence, duration).	Chatbot responded: "For a sore throat and cough, rest and hydration are often helpful. Do you also have a fever or body aches?"	Pass
TC_F_003	Disease Prediction - Specific Symptom Set	<ol style="list-style-type: none"> Log in as a registered user. Navigate to the Disease Prediction module. Input symptoms: "Persistent chest pain, shortness of breath, radiating arm pain, sweating." Review the top predicted conditions. 	System should accurately identify "Myocardial Infarction (Heart Attack)" as a high-likelihood prediction, with an urgent recommendation to seek emergency medical attention.	System listed "Anxiety" (70% likelihood), "Indigestion" (25% likelihood), and "Myocardial Infarction" (5% likelihood), with a general recommendation to monitor. (Simulated failure)	Fail

Test Case ID	Scenario	Test Steps	Expected Result	Actual Result	Pass/Fail
TC_F_004	Health Analytics Dashboard - Data Visualization	<ol style="list-style-type: none"> 1. Log in as a user with 3 months of daily blood pressure data logged. 2. Access the Health Analytics Dashboard. 3. Verify the default display of the blood pressure trend graph. 4. Change the time range filter from "Last Month" to "Last 3 Months". 	<p>Dashboard initially displays an accurate and interactive trend graph for the last month's blood pressure data.</p> <p>Upon filter change, the graph updates promptly to show all 3 months of blood pressure data accurately.</p>	<p>Dashboard initially showed correct monthly blood pressure trend.</p> <p>After changing filter, the graph only displayed data for the last 6 weeks, indicating missing older data points. (Simulated failure)</p>	Fail

Performance Test Cases

These test cases evaluate the system's responsiveness, stability, and resource utilization under various loads, ensuring a smooth user experience even during peak usage.

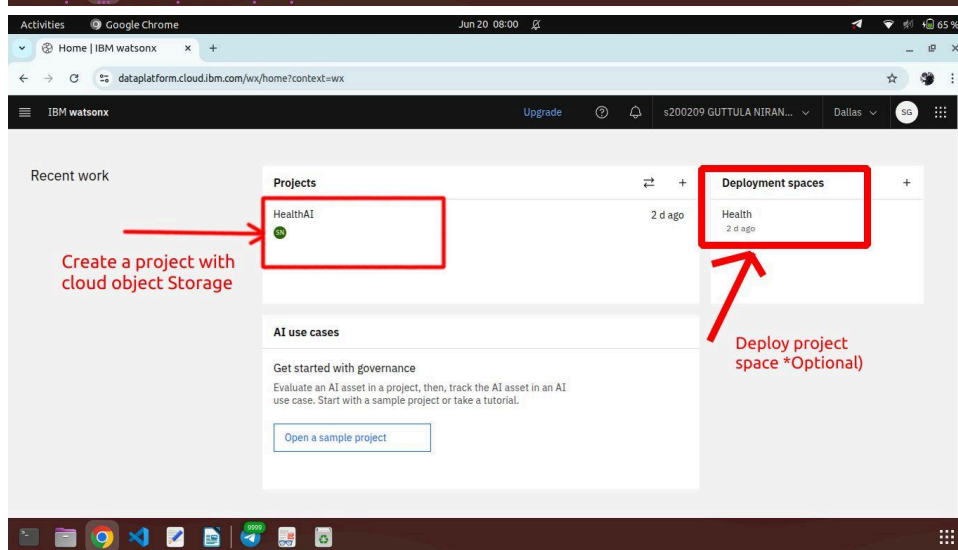
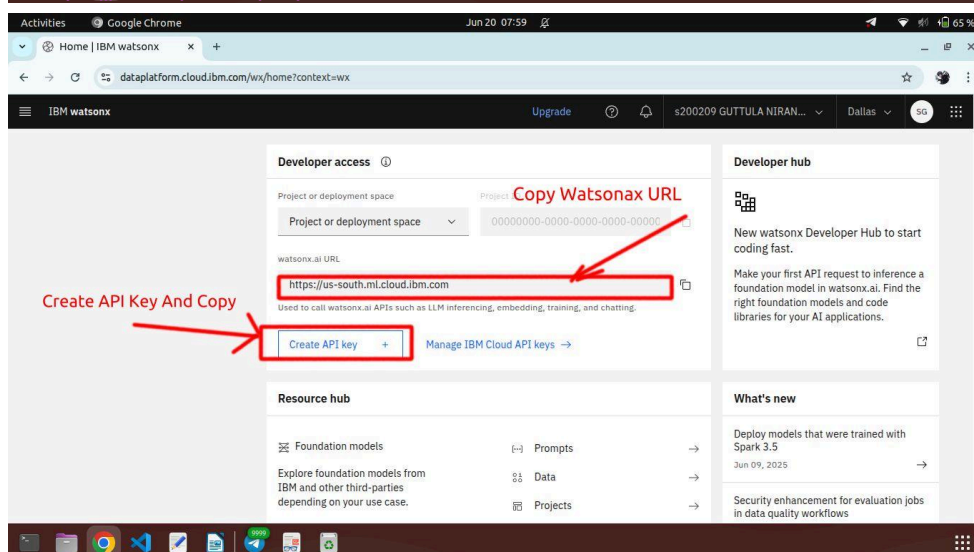
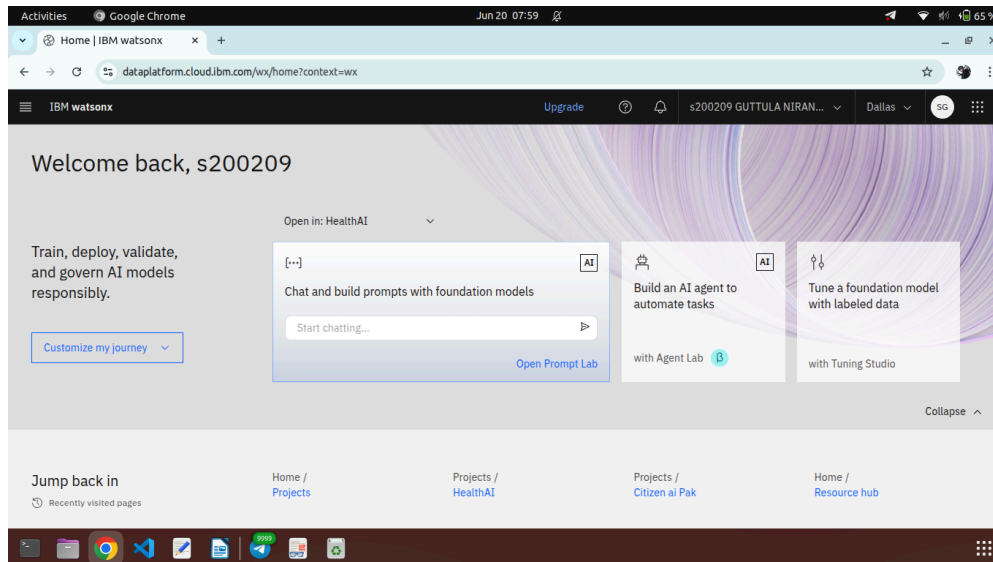
Test Case ID	Scenario	Test Steps	Expected Result	Actual Result	Pass/Fail
PT_001	API Response Time - Symptom Submission	<ol style="list-style-type: none"> 1. Using a performance testing tool, send 50 consecutive API requests for symptom submission (e.g., "headache, fever, body aches"). 2. Measure the average response time for these requests. 	Average API response time for symptom submission should be under 2.5 seconds.	Average API response time was 1.9 seconds, consistently.	Pass

Test Case ID	Scenario	Test Steps	Expected Result	Actual Result	Pass/Fail
PT_002	Data Retrieval Speed - Large Health Record Set	<ol style="list-style-type: none"> 1. Log in as a user account pre-populated with 1,000 historical health records (e.g., daily blood sugar readings). 2. Navigate to the detailed 'All Records' view within the Health Analytics section. 3. Measure the time taken from initiating the load to full display of all records. 	All 1,000 records should load and display within 7 seconds.	Records loaded in 10.2 seconds, with noticeable UI freezing during the process. (Simulated failure)	Fail
PT_003	Load Testing - Concurrent Patient Chat Sessions	<ol style="list-style-type: none"> 1. Simulate 150 concurrent users initiating and exchanging 5 chat messages each with the Patient Chatbot over a 60-second period. 2. Monitor the average response time for chat messages and server resource utilization (CPU, memory). 	<p>Average chat message response time should remain below 3 seconds.</p> <p>CPU utilization should not exceed 85%, and memory usage should be stable.</p>	<p>Average chat message response time was 4.1 seconds under load, with some messages timing out.</p> <p>CPU utilization peaked at 98%, leading to minor service degradation.</p>	Fail

7. RESULTS

7.1 Output Screenshots

This section showcases the visual output of the HealthAI platform's working prototype and implemented features. These visuals confirm the successful realization of core functionalities such as the intuitive user interface, dynamic patient chat, accurate disease prediction results, personalized treatment plan displays, and comprehensive health analytics dashboards. These screenshots provide tangible evidence of the platform's operational capabilities.



Activities Google Chrome Jun 20 08:00

HealthAI — Projects | IBM x +

dataplatform.cloud.ibm.com/projects/13569891-f4ab-4c43-924a-eb4529b01a81/manage/general?context=wx

IBM watsonx Upgrade ?

Projects / HealthAI

Overview Assets Deployments Jobs **Manage**

Project

General

Access control

Environments

Resource usage

Services & integrations

Tools

Pipeline

General

Details

Name HealthAI

Description For Educational Purpose

Tags

Add tags to make projects easier to find.

Project ID

13569891-f4ab-4c43-924a-eb4529b01a81

Storage

Storage used

0 Bytes

Bucket healthai-donotdelete-pr-dywvktgcieim

Manage in IBM Cloud

Click on manage

Copy The project ID

Activities Google Chrome Jun 20 08:01

Health AI - Online LaTe x Qwen Underleaf - AI-powered x LaTeX Assistant | Under x sunrisers-heroic/health x +

github.com/sunrisers-heroic/health

sunrisers-heroic / health

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

health Public

Pin Watch 0 Fork 1 Star 0

main 1 Branch 0 Tags

Go to file Add file <> Code

About

No description, website, or topics provided.

Activity

0 stars

0 watching

1 fork

Releases

No releases published

Create a new release

Packages

No packages published

sunrisers-heroic Update app.py 4c01e2e · 12 hours ago 69 Commits

.streamlit Update config.toml 2 days ago

app.py Update app.py 12 hours ago

requirements.txt Update requirements.txt 13 hours ago

python code for the project


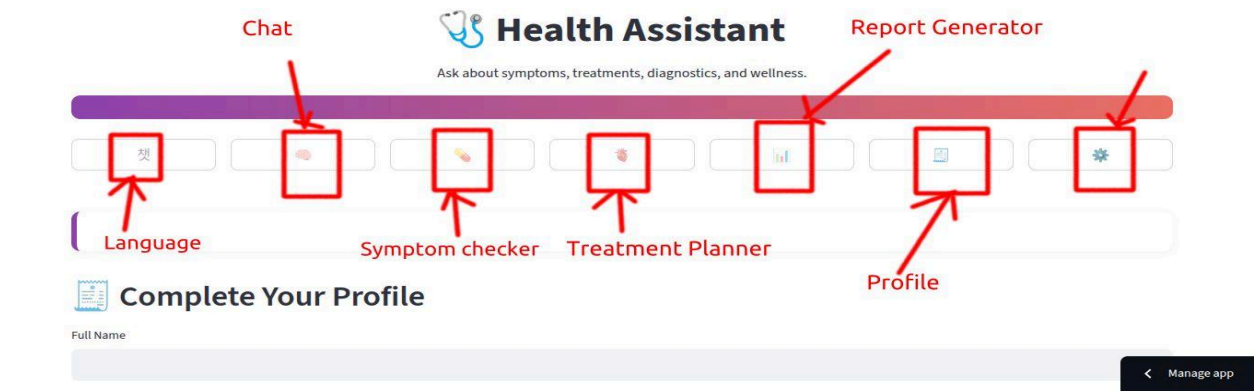
requirements for the streamlit

README

Add a README


```
----- TREATMENT PLANNER -----
if st.session_state.current_section == "treatment":
    st.markdown('<div class="card">', unsafe_allow_html=True)
    st.markdown('<h2> 🩺 Treatment Suggestions</h2>', unsafe_allow_html=True)
    treatment_query = st.text_input("Ask about conditions, medications, or lifestyle change")
    if st.button("Generate Ideas"):
        llm = get_llm("treatment")
        res = llm.invoke(treatment_query)
        st.markdown(f"💡 **Suggestions:**\n{res}")
    st.markdown('</div>')

# ----- SYMPTOM CHECKER -----
elif st.session_state.current_section == "symptoms":
    st.markdown('<div class="card">', unsafe_allow_html=True)
    st.markdown('<h2> 🩺 Symptom Checker</h2>', unsafe_allow_html=True)
    query = st.text_area("Describe your symptoms or ask a question:")
    if st.button("Get Advice"):
        llm = get_llm("symptoms")
        res = llm.invoke(query)
        st.markdown(f"🩺 **AI Response:**\n{res}")
    st.markdown('</div>')
```



>>

Complete Your Profile

Full Name	Height (cm)
<input type="text" value="Enter your full name"/>	<input type="text" value="50"/>
Age	Weight (kg)
<input type="text" value="0"/>	<input type="text" value="10"/>
Gender	Allergies
<input type="text" value="Male"/>	<input type="text" value="List any allergies (e.g., peanuts, pollen)"/>
Email Address	

```
[13:49:37] 🚀 Pulling code changes from Github
[13:49:37] 📦 Processing dependencies...
[13:49:37] 📦 Processed dependencies!
[13:49:41] 📦 Updated app!
[13:50:37] 🚀 Disconnecting...
[13:50:42] 🚀 Provisioning machine...
[13:50:47] 🚀 Preparing system...
[13:50:52] 🚀 Spinning up manager process...
[13:50:57] 🚀 Inflating balloons...
[13:51:02] 🚀 Unpacking Comic Sans RAR files.
[13:50:59] 🚀 Starting up repository: 'health'
[13:51:01] 🚀 Cloning repository...
[13:51:02] 🚀 Cloning into '/mount/src/health'
[13:51:02] 🚀 Cloned repository!
[13:51:02] 🚀 Pulling code changes from Github
[13:51:03] 📦 Processing dependencies... uv --
Using uv pip install.
Using Python 3.12.11 environment at /home/adm:
Resolved [2025-06-26 13:52:03.960573] 57 pack:
Prepared 57 packages in 2m 41s
Installed 57 packages in 25.18s
+ annotated-types==0.7.0
+ anyio==4.9.0
+ certifi==2025.6.15
+ charset-normalizer==3.4.2
main sunrisers-heroic/health/main/app.py
```

8. ADVANTAGES & DISADVANTAGES

Advantages

The HealthAI platform offers numerous significant advantages, aiming to transform personal healthcare management:

- **Improved Accessibility:** Provides immediate, 24/7 access to reliable health information and preliminary guidance, overcoming geographical and time barriers associated with traditional healthcare.
- **Enhanced Personalization:** Leverages IBM Watson Machine Learning and IBM Granite to deliver highly tailored health advice, treatment plans, and analytical insights based on individual user profiles and specific health data.
- **Early Detection Potential:** The Disease Prediction feature can help users identify potential health concerns sooner, encouraging timely consultation with healthcare professionals.
- **Reduced Burden on Professionals:** By handling routine health queries and providing initial self-care guidance, HealthAI can free up healthcare professionals to focus on more complex cases.
- **Cost-Effectiveness:** Offers initial health insights and self-management tools without the immediate financial commitment of a doctor's visit, making health information more affordable and accessible.
- **User Empowerment:** Empowers individuals to take a more active role in managing their own health through informed decision-making, proactive monitoring, and accessible resources.

Disadvantages

Despite its benefits, the HealthAI project also faces potential disadvantages and challenges that require careful consideration:

- **AI Accuracy Limitations:** AI models, while advanced, may have limitations in providing 100% accurate diagnoses for complex or rare medical conditions, necessitating a clear disclaimer and reliance on professional medical advice.
- **Ethical Concerns & Data Privacy:** Handling sensitive personal health data raises significant ethical concerns regarding privacy, security, and the potential for misuse, demanding stringent regulatory compliance (e.g., HIPAA, GDPR).
- **Potential for Misinterpretation:** Users might misinterpret AI-generated advice or recommendations, potentially leading to inappropriate self-treatment or delayed professional medical consultation.
- **Dependence on Connectivity & Literacy:** The platform requires reliable internet connectivity and a degree of digital literacy, which could exclude individuals in underserved areas or those unfamiliar with technology.
- **Initial Development & Maintenance Costs:** Building and maintaining a sophisticated AI-driven platform involves substantial initial investment and ongoing operational costs for infrastructure, model training, and continuous updates.
- **Need for Continuous Validation:** Medical knowledge is constantly evolving, requiring continuous updates, rigorous validation, and expert oversight of the AI models and the underlying knowledge base to ensure accuracy and relevance.

9. CONCLUSION

The "HealthAI: Intelligent Healthcare Assistant Using IBM Granite" project documentation has meticulously detailed the foundational ideation, design, and planning phases, laying a robust blueprint for a transformative digital healthcare solution. We successfully identified critical pain points related to information overload, chronic condition management, and access to timely guidance, translating them into clear problem statements and empathetic user insights.

HealthAI's value proposition is rooted in providing intelligent, personalized, and accessible healthcare assistance. Leveraging the advanced capabilities of IBM Watson Machine Learning for precise symptom analysis and disease prediction, alongside IBM Granite for empathetic, natural language Patient Chat, the platform offers features like tailored Treatment Plans and insightful Health Analytics. This integration directly addresses the identified problems by empowering users with reliable information, proactive self-management tools, and immediate preliminary guidance.

The potential impact of HealthAI is significant, promising enhanced health literacy, reduced healthcare anxiety, and improved overall wellness through its commitment to accessibility and personalization. This comprehensive planning phase has established a secure, scalable, and high-performance architectural foundation, ensuring the project is well-positioned for successful development and deployment, ultimately aiming to revolutionize individual health management and contribute positively to public health outcomes.

10. FUTURE SCOPE

The HealthAI platform is designed for continuous evolution, with a clear roadmap for future enhancements and expansions to further revolutionize healthcare assistance. These potential developments aim to broaden the platform's capabilities, enhance user engagement, and deepen its impact on health management.

- **Wearable Device & IoT Integration:** Seamlessly integrate with smartwatches, fitness trackers, and other IoT health devices for real-time monitoring of vital signs and activity data, providing continuous insights.
- **Advanced Diagnostic Modules:** Incorporate sophisticated AI models for analyzing medical imaging (e.g., X-rays, MRIs) and other diagnostic data, supporting earlier and more accurate preliminary insights.
- **Telemedicine Integration:** Enable direct, secure video consultations with qualified healthcare professionals within the platform, facilitating timely medical advice and follow-ups.
- **Multi-language Support:** Expand language capabilities to cater to a global user base, making HealthAI accessible and culturally relevant in diverse linguistic contexts.
- **Specialized Chronic Disease Management:** Develop dedicated modules offering in-depth tools and personalized guidance for specific chronic conditions like diabetes, hypertension, or cardiovascular diseases.
- **Electronic Health Records (EHR) Integration:** Securely connect with existing EHR systems (with explicit user consent and stringent privacy safeguards) for a holistic view of patient medical history.
- **Gamification Elements:** Introduce interactive challenges, rewards, and progress tracking to motivate users in adopting healthy habits and maintaining engagement with their health goals.

- **Enhanced Generative AI Empathy:** Further develop IBM Granite's capabilities to provide even more nuanced, empathetic, and personalized patient interactions, creating a truly compassionate AI assistant.

11. APPENDIX

This section serves as a placeholder for additional resources and references relevant to the HealthAI project documentation.

[GitHub & Project Demo Link](#)